

1 Agents and Abstraction

Definition 1.1 (Planning Horizon):

- **Static:** The world does not change over time.
- **Finite Horizon:** The agent reasons about a fixed finite number of time steps. (Agent know when it will end)
- **Indefinite Horizon:** The agent reasons about a finite but not predetermined number of time steps, such as until goal completion. (Agent know it will end, but not sure when)
- **Infinite Horizon:** The agent plans as if it will continue operating forever. (agent know it will never end)

Definition 1.2 (Representation):

- **Explicit States:** A state represents one possible configuration of the world.
- **Features:** Natural descriptors of states. (binary features can represent exponentially many states)
- **Individuals and Relations:** Use feature for reasoning about individuals and their relationships without necessarily knowing all individuals or when there are infinitely many individuals.

Definition 1.3 (Computational Limits):

- **Perfect Rationality:** The agent always selects the optimal action, which is often not possible in practice.
- **Bounded Rationality:** The agent selects a possibly sub-optimal action given its limited computational resources.

Definition 1.4 (Uncertainty):

- **fully observable:** agent knows the state of the world from the observation
- **partially observable:** there can be many state that are possible given an observation.

Definition 1.5 (Uncertain Dynamics):

- **Deterministic:** The outcome of an action is always the same.
- **Stochastic:** There is uncertainty over the states resulting from executing a given action.

Definition 1.6 (Goals or Complex Preferences):

- **Achievement Goals:** Goals that an agent aims to achieve, which can be represented as complex logical formulas.
- **Maintenance Goals:** Goals that an agent seeks to maintain over time.
- **Complex Preferences:** Involves trade-offs between various desiderata, potentially at different times, and may be either ordinal or cardinal (e.g: medical).

Definition 1.7 (Reasoning by Number of Agents):

- **Single Agent Reasoning:** The agent assumes any other agents are part of the environment, focusing on individual goal achievement.
- **Adversarial Reasoning:** The agent considers another agent acting in opposition to its goals, common in competitive settings.
- **Multi-agent Reasoning:** The agent strategically reasons about the actions and goals of other agents, which may be cooperative, competitive, or independent.

Algorithm	Frontier	Runtime	Space	halts?
uninformed (no heuristic)				
Depth-First Search	LIFO	Exp / $O(b^m)$	Linear / $O(bm)$	No
Breadth-First Search	FIFO	Exp / $O(b^d)$	Exp / $O(b^d)$	Yes
Lowest-Cost-First Search	Lowest cost	Exp	Exp	Yes
Dijkstra's Algorithm*	Lowest cost	$O((V + E) \log V)$	$O(V^2)$	
Iterative-Deepening Search*	LIFO in FIFO ¹	Exp / $O(b^d)$	Linear / $O(bd)$	Yes ²
informed (has heuristic)				
(Greedy) Best-First Search	Global min heuristic	Exp	Exp	No
Heuristic Depth-First Search	Local min heuristic (LIFO)	Exp / $O(b^m)$	Linear / $O(bm)$	No
A* Search	Lowest (cost + heuristic)	Exp	Exp	Yes

b is the branching factor

m is the maximum depth of the search tree

d is the depth of the shallowest goal node.

¹ a BFS but for every depth limit do a DFS

² Guaranteed to terminate at depth d

Algorithm	Completeness	Optimality
uninformed (no heuristic)		
Depth-First Search	No (fails for infinite cycle)	No (not considering all possibilities)
Breadth-First Search	Yes	Yes (if cost is uniform)
Lowest-Cost-First Search	Yes	Yes
Dijkstra's Algorithm*	Yes	Yes
Iterative-Deepening Search*	Yes (Same as BFS)	No (but guaranteed shallowest goal)
informed (has heuristic)		
(Greedy) Best-First Search	No (fails for infinite cycle)	No (from not considering cost of arc)
Heuristic Depth-First Search	No (fails for infinite cycle)	No (not considering all possibilities)
A* Search	Yes ¹	Yes ¹

¹ Assuming heuristic is **admissible**, branch factor is **finite**, and arc cost are bounded **above zero**

If h satisfies the **monotone restriction**, A* with **multiple path pruning** always finds the **shortest path** to a goal.

2 Graph Search Algorithm

Definition 2.1 (Admissible) An **Admissible** heuristic **never overestimate** the cost from **any node to the goal**. An **Admissible** search algorithm returns an **optimal solution** if it exists.

Definition 2.2 (Monotone / Consistent) A heuristic function h satisfies the **monotone restriction** if $h(m) - h(n) \leq \text{cost}(m, n)$ **for every arc** $\langle m, n \rangle$. (The heuristic of a path is always less than or equal to the true cost)

Monotonicity is like **admissibility** but between any two nodes.

So, a **consistent** heuristic is **admissible**, but a **admissible** heuristic is not necessarily **consistent**.

3 Adversarial Search(Minimax)

- function minimax (node, depth, maximizingPlayer) is


```

      if depth = 0 or node is a terminal node then
        return the heuristic value of node
      if maximizingPlayer then
        value := -inf
        for each child of node do
          value := max(value, minimax(child, depth - 1, FALSE))
        return value
      else (* minimizing player *)
        value := inf
        for each child of node do
          value := min(value, minimax(child, depth - 1, TRUE))
        return value
      
```
- **Suitable Type of Problem:**
 - Competitive two-person, zero-sum games.
 - Two players take turns to move and the one winner one loser.
- **Idea:**
 - Find **best option** for you on nodes you control (MAX)
 - Assumes opponent will take **worst option** for you on their node (MIN)
 - Recursively search leaf nodes and percolate optimal value upward
- **Pruning Methods:**
 - **Alpha-beta Pruning:**
 - * Ignore portions of the search tree without losing optimality
 - * doesn't change worst-case performance (Exp)
 - **Heuristic Pruning (Early Stopping):**
 - * Heuristics are used to evaluate the potential of non-terminal states.
 - * This method saves computational resources but may not always yield the optimal solution.

4 Higher level strategies

Search	Search Complexity	Difficulty	Reason to Win
Symmetric	b^n	Not able to construct backward on dynamically constructed graph	Choose between forward / backward search based on branching factor
Bidirectional	$2b^{\frac{k}{2}} \ll b^k$	Make sure frontiers meet	Searches forward and backward simultaneously, leading to exponential savings in time and space.
Island-Driven	$mbk^{\frac{k}{m}} \ll b^k$	identify islands hard to guarantee optimality	Decomposes the problem into m smaller subproblems, each of which is easier to solve.

5 Constraint Satisfaction Problems (CSPs)

- **Definition:**

- Set of **variables**, **domain for each variable**, set of **constraints** or **evaluation function**.
- **Solution** is an assignment to the **variables** that satisfies **all constraints**.
- **Solution** is a **model of the constraints**.

- **Problem Types:**

- **Satisfiability Problems:** Find assignment satisfies the given hard constraints.
- **Optimization Problems:** Find assignment optimizes the evaluation function(soft constraints).

- **Search Representation:**

Assignment Type	Description
Complete Assignment	Node is assignment of value to all variables. Neighbours are created by changing one variable value.
Partial Assignment	Nodes is assignment to the first $k - 1$ variables. Neighbours are formed by assigning a value to the k^{th} variable.

- **Dual Representations of Crossword Puzzle:**

Type	Nodes	Domains	Constraints
Primal	word positions	letters	intersecting letters are same
Dual	squares	letters	words must fit

- **Example of CSP Setup:**

Problem	Variables	Domains	Constraints
Crosswords	letters	a-z	words in dictionary
Crosswords	words	dictionary	letters match
Scheduling	times events resources	times,dates types values	before, after same resource
Party Planning	guests	values	cliques
Ride Sharing	people/trips	locations	cars

- **Constraints & Solution**

- **Constraints:** Can be **N-ary** or **Binary**.
- **Solutions:**

- **Generate and Test**

Exhaustively check all combinations against constraints.

- **Backtracking**

Prune large portions of the state space by ordering variables and evaluating constraints.

Efficiency depends on **order** of variables.

Find optimal ordering is **as hard** as solving the problem.


Cut off large branches as soon as possible, push failures as high as possible.

- **Consistency Techniques**

Look for inconsistencies to simplify the problem / Graphical representation

- **Constraint Network (CN)**

- **Domain constraint:**
unary constraint of values x on values in a domain, $\langle X, c(X) \rangle$
- **Domain consistent:**
A **node** is **Domain consistent** if no domain value violates any domain constraints.
A CN is **Domain consistent** if all nodes are **Domain consistent**.
- **Arc** $\langle X, c(X, Y) \rangle$ is:
A **constraint** on X posed by Y .
Arc consistent if for all $X \in D_X$, there exist some $Y \in D_Y$ such that $c(X, Y)$ is satisfied.
- CN is **Arc consistent** if all arcs are **arc consistent**.
- set of variables $\{X_1, X_2, \dots, X_N\}$ is **path consistent** if all **arcs** and **domains** are **consistent**



- **AC-3 (CN) algorithm** (Alan Mackworth, 1977)

- **Purpose:** Makes a Constraint Network (CN) arc consistent and domain consistent.
- **Procedure:**
 - * Initialize the To-Do Arcs Queue (TDA) with all inconsistent arcs.
 - * Make all domains domain consistent.
 - * Put all arcs in TDA.
 - * Repeat until TDA is empty:
 - Select and remove an arc $\langle X, c(X, Y) \rangle$ from TDA.
 - Remove all values from the domain of X that:
 - do not have a corresponding value in the domain of Y
 - satisfying the constraint $c(X, Y)$.
 - If any values were removed, for all $Z \neq Y$, add back arcs $\langle Z, c'(Z, X) \rangle$ into TDA .
(Add back all constraints posed to other variable by X. As the X value enforced by the constraints / arc we removed is used by some other constraints posted by X to other variables)
- **Termination:**
 - * AC-3 always terminates under one of three conditions:
 - **Every domain is empty:** no solution.
 - **Every domain has a single value:** a solution.
 - **Some domains have 1+ value:** not sure if a solution exists.
(further splitting and recursive call needed)
- **Properties:**
 - * Termination is guaranteed.
 - * Time complexity is $O(cd^3)$.¹
 - * Consistency of each arc can be checked in $O(d^2)$ time.
- **Different elimination ordering** can result in **different size** of intermediate constraints.

¹where n is the number of variables, c is the number of binary constraints, and d is the maximum size of any domain

– Variable Elimination

* Concept:

- Variables are eliminated one by one, transferring their constraints to neighbours.
- A single remaining variable with no values indicates an **inconsistent** network.

* Algorithm:

- If only one variable remains, return the intersection of the unary constraints involving it.
- Select a variable X .

Join the constraints where X appears to form a new constraint R .

Project R onto other variables to form R_2 .


Place new constraint R_2 between all variables previously connected to X .

Remove X from the problem.

Recursively solve the simplified problem.

Return R joined with the solution from the recursive call.

- * Finding the optimal elimination ordering is as complex as the CSP itself.



• Local Search: (Back to CSP as Search)

- Maintain a variable assignment, select neighbors to improve heuristic value, and stop when a satisfying assignment is found, or return the best assignment found.
- Aim is to find an assignment with zero **unsatisfied constraints (Conflict)**
- Goal is an assignment with **zero conflicts** (heuristic: # of conflicts)

• Greedy Descent Variants:

- At every step:
 - (Select the variable-value pair that **minimize** # of conflicts)
 - (Select a variable involved in the **most** # of conflicts, then a value **minimize** # of conflicts)
 - (Select a variable involved in **any** conflicts, then a value **minimize** # of conflicts)
 - (Select a variable **at random**, then a value **minimize** # of conflicts)
 - (Select a variable and value **at random**, accept if **doesn't increase**² # of conflicts)

²Sometime accept even increase # of conflicts to escape local minimum

- **GSAT (Greedy SATisfiability):**

- Start with a random assignment of values to all variables n
heuristic $h(n) = \#$ of unsatisfied constraints
- repeat until heuristic becomes 0 (Solved):
Evaluate neighbours of n (not n , cannot change same variable twice in a row);
Let n be the neighbour n' that minimizes the heuristic, even if $h(n') > h(n)$.

- **Problems:**

stuck at local minimum, cannot pass a **plateau** where $h(n)$ are uninformative.

Ridge is a local minimum where **n-step look-ahead** might help.


- **Randomized GSAT:** allow move to a random neighbour, reassign all variable randomly.

- **Tabu lists:**

Maintain a tabu list of k last assignments to prevent cycling.


Reject assignments exist on tabu lists.³

More efficient than a list of complete assignments, but expensive if k is large.



- **Stochastic local search** is a mix of:

- **Greedy descent:** Move to a lowest neighbour (GSAT)
- **Random walk:** taking some random steps
- **Random restart:** reassigning all values randomly



- **Simulated Annealing** (Variant of Stochastic local research)

- Pick a random value for a random variable (neighbour):
Adopt if its an improvement.
If it's not an improvement, adopt it probabilistically based on temperature T . (High \rightarrow Low)
(move from current assignment n to new assignment n' with probability $e^{-\frac{h(n')-h(n)}{T}}$)⁴
Terminate when criteria is met.

- **Parallel Search:**

- Maintains a population of k individuals (total assignments).
- Updates each individual in the population at every stage.
- Reports when any individual is a solution.
- Operates like k restarts but with k times the minimum steps.

- **Beam Search:**

- Similar to parallel search with k individuals, but **choosing the k best from all neighbours**.
(All if there are less than k)
- Reduces to greedy descent when $k = 1$.
- The value of k limits space and facilitates parallelism.

³e.g: $k = 1$ means reject assignment of the same value to the variable chosen.

⁴difference in heuristic value divided by temperature

- **Stochastic Beam Search:**

- A probabilistic variant of beam search, choosing k individuals for the next generation.
(probability of a neighbour n is chosen is proportional to $e^{-\frac{h(n)}{T}}$)
- Maintains diversity among individuals and reflects their fitness (heuristic).
- Operates like asexual reproduction, with mutation allowing fittest individuals to prevail.

- **Genetic Algorithms:**

- Like stochastic beam search but **combines pairs of individuals to create offspring**.
- Fittest individuals are more likely to be chosen for reproduction.
- Crossover and mutation (change some value) to form new solutions.
- Continues until a solution is found.

- **Crossover:**

- Given two individuals, each offspring's attributes are randomly chosen from one of the parents.
- The effectiveness depends on the ordering of variables, many variations are possible.

- **Comparing Stochastic Algorithms:**

- Considers how to compare algorithms with different success rates and runtimes.
- Acknowledges that summary statistics like mean, median, or mode runtimes may not be informative.

- **Runtime Distribution:**

- Plots runtime or steps against the proportion of runs solved within that time.
- Helps in understanding the performance distribution of stochastic algorithms.

6 Inference and Planning

-